

Variational Inference

Introduction: Bayesian networks

Helge Langseth and Thomas Dyhre Nielsen*

Oct. 2018

* Some of the introductory slides are stolen from Manfred Jaeger, Aalborg University.



Universidad de Almería:

- Antonio Salmerón
- Andrés Masegosa
- Rafael Cabañas de Paz
- Rafael Rumí
- Darío Ramón-Lopez

Aalborg University:

- Thomas D. Nielsen
- Anders L. Madsen


NTNU:

- Helge Langseth



Simultaneous Localization and Mapping: learn a map of the environment and locate current position

Example 2: Image Segmentation



(source: <http://pubs.niaaa.nih.gov/publications/arh313/243-246.htm>)


- Divide image into small number of regions representing structurally similar areas

Given a collection of texts:



Goal: automatically learn semantic descriptors for documents and words, that support document clustering, text understanding, information retrieval ...

Micro-array gene expression data:



Which genes are expressed under which conditions? Which are co-regulated, or functionally dependent?

Common ground in 4 examples:

- Use a **probabilistic model**, typically learned from data (using **statistical learning techniques**)
- Apply probabilistic inference algorithms to use models for prediction (**classification, regression**), structure analysis (**clustering, segmentation**)

Advantages of probabilistic/statistical methods:

- Principled quantification of prediction uncertainties
- Robust and principled techniques to deal with incomplete information, missing data.

Need: probabilistic models that ...

- can represent models for high-dimensional state spaces
- support efficient learning and inference techniques

Probabilistic Graphical Models ...

- support a structured specification of high-dimensional distributions in terms of low-dimensional factors
- structured representation can be exploited for efficient learning and inference algorithms (sometimes ...)
- graphical representation gives human-friendly design and description possibilities

This seminar series is about efficient inference and learning for Bayesian networks using “Variational inference”.

Day 1: Bayesian networks – Definition and inference

- Definition of Bayesian networks: Syntax and semantics
- Exact inference
- Approximate inference using MCMC

Day 2: Variational inference – Introduction and basis

- Approximate inference through the *Kullback-Leibler divergence*
- *Variational Bayes*
- The *mean-field* approach to Variational Bayes

Day 3: Variational Bayes – cont'd

- Solving the VB equations
- Introducing Exponential families

Day 4: Scalable Variational Bayes

- Variational message passing
- Stochastic gradient ascent
- Stochastic variational inference

Day 5: Current approaches and extensions

- Variational Auto Encoders
- Black Box variational inference
- Probabilistic Programming Languages

Online repository

Materials for the course, including (almost last-version of) the slides, and description of implementation-tasks are available at

<https://github.com/HelgeLangseth/AlmeriaVB>

The screenshot shows the GitHub repository page for HelgeLangseth/AlmeriaVB. At the top, the repository name is displayed with navigation links for Code, Issues (0), Pull requests (0), Projects (0), Wiki, Insights, and Settings. On the right, there are buttons for Unwatch (2), Star (2), and Fork (0). Below the navigation bar, the repository description "Slides and Jupyter notebooks for VB course in Almeria" is shown with an "Edit" button. A "Manage topics" link is also present. A summary bar indicates 7 commits, 1 branch, 0 releases, 1 contributor, and the Apache-2.0 license. Below this, there are buttons for "Branch: master", "New pull request", "Create new file", "Upload files", "Find file", and "Clone or download". A commit history table follows, listing recent commits by HelgeLangseth, including additions to Python, Slides, LICENSE, and README.md. The bottom section shows the README.md file content, which includes the title "AlmeriaVB" and a description of the repository's purpose and file naming convention.

HelgeLangseth / AlmeriaVB

Unwatch 2 Star 2 Fork 0

Code Issues 0 Pull requests 0 Projects 0 Wiki Insights Settings

Slides and Jupyter notebooks for VB course in Almeria Edit

Manage topics

7 commits 1 branch 0 releases 1 contributor Apache-2.0

Branch: master New pull request Create new file Upload files Find file Clone or download

File	Commit Message	Time Ago
Python	Added Thomas' MCMC exercise	2 days ago
Slides	Added Thomas' MCMC exercise	2 days ago
LICENSE	Initial commit	4 days ago
README.md	Update README.md	4 days ago

README.md

AlmeriaVB

This repository contains slides and Jupyter notebooks for VB course in Almeria in Oct 2018.

The slides (in the `slides` folder) are named with the lecture number (1 to 5).

I will assume that these topics are (well) known:

- Probabilities, $P(X = x)$; conditional probabilities, $P(X = x | Y = y)$.
- Independence, $\mathbf{X} \perp\!\!\!\perp \mathbf{Y}$; conditional independence $\mathbf{X} \perp\!\!\!\perp \mathbf{Y} | \mathbf{Z}$.
- “Standard” probability calculus:

Product: $P(x, y) = P(x | y) \cdot P(y) = P(x) \cdot P(y | x)$.

Sum-rule: $P(x \vee y) = P(x) + P(y) - P(x \wedge y)$.

Total probability: $P(Y = y) = \sum_{x'} P(y | X = x') \cdot P(X = x')$.

Bayes rule: $P(x | y) = P(x) \cdot P(y | x) / P(y)$.


I will assume that these topics are (somewhat) known:

- Bayesian network syntax and semantics.
- Exact inference in Bayesian networks.
- Approximate inference using MCMC.

Implementation tasks:


- We will get some implementation tasks as we move along.
- You will be supplied partly running Python-code (in the form of Jupyter notebooks).
- You will need to have **Python 3.x** on your computer (<https://www.python.org/downloads/>), and a set of packages (numpy, scipy, matplotlib, jupyter).

A simple example: “Explosion”




$$P(E, L, G, X, C)$$

A simple example: “Explosion”




$$P(E, L, G, X, C)$$

A simple example: “Explosion”



$$P(E, L, G, X, C)$$

A simple example: “Explosion”



$$\text{pa}(X) = \{L, G\}$$


$$\text{nd}(X) = \{E, L, G\}$$

$$X \perp\!\!\!\perp E \mid \{L, G\}$$

Other d-sep. rules: Jensen&Nielsen (07)

$$P(E, L, G, X, C)$$


A simple example: “Explosion”



Other d-sep. rules: Jensen&Nielsen (07)

$$\begin{aligned}
 P(E, L, G, X, C) &= P(E) \cdot P(L \mid E) \cdot P(G \mid E, L) \cdot P(X \mid E, L, G) \cdot P(C \mid E, L, G, X) \\
 &= P(E) \cdot P(L \mid E) \cdot P(G \mid E) \cdot P(X \mid L, G) \cdot P(C \mid X)
 \end{aligned}$$

Markov properties \Leftrightarrow Factorisation property



In the distribution P defined by the BN the following independence relation holds:

$$P(X_i \mid pa(X_i), rest(X_i)) = P(X_i \mid pa(X_i))$$

“ X_i is independent of its non-descendants given its parents”

“rest” are all nodes except:

- X_i itself
- X_i 's parents
- X_i 's descendants

Bayesian network syntax

Let X_1, \dots, X_n be a collection of random variables. A Bayesian network over X_1, \dots, X_n consists of

- a directed acyclic graph $\mathcal{G} = \{\mathcal{V}, \mathcal{E}\}$ whose nodes \mathcal{V} are the variables X_1, \dots, X_n
- a set of local conditional distributions, $\mathcal{P} = \{p(X_i \mid \text{pa}(X_i)), X_i \in \mathcal{V}\}$, where $\text{pa}(X_i)$ are the parents of X_i in \mathcal{G} as defined by the edges \mathcal{E} .

Bayesian network semantics

A Bayesian network \mathcal{N} with nodes X_1, \dots, X_n defines a joint distribution

$$p_{\mathcal{N}}(X_1, \dots, X_n) = \prod_{i=1}^n p(X_i \mid \text{pa}(X_i))$$

Exact inference

Bayesian network inference

Inference in the Bayesian network amounts to calculating $p(\mathbf{Z} = \mathbf{z} \mid \mathbf{X} = \mathbf{x})$, where


- $\mathbf{X} \subset \mathcal{V}$ are the observed variables, currently taking the configuration $\mathbf{X} = \mathbf{x}$.
- $\mathbf{Z} \subseteq \mathcal{V} \setminus \mathbf{X}$ are our variables of interest.

Inference is therefore the tool to answer any probabilistic query we may have in our domain, given a partial (or empty) observation from the domain

- *“What is the chance that Almería will move back to La Liga before 2020, given that they didn’t win during the first four games of this campaign?”*
- *“What is the probability of Google going bankrupt before you are done with your education?”*

Probability propagation

$$\left. \begin{array}{l} \text{Bayesian network} \\ + \\ \text{Evidence: } \mathbf{X}_E = \mathbf{x}_e \end{array} \right\} \Rightarrow P(\mathbf{X}_Q | \mathbf{x}_e)?$$




Example: Calculate $P(X = x | G = g)$

$$\begin{aligned} P(x, g) &= \sum_e \sum_l \sum_c P(E = e, L = l, g, X = x, C = c) \\ &= \sum_e \sum_l \sum_c P(e) \cdot P(l | e) \cdot P(g | e) \cdot P(x | l, g) \cdot P(c | x) \\ &= \sum_e P(e) \cdot P(g | e) \sum_l P(l | e) \cdot P(x | l, g) \sum_c P(c | x) \end{aligned}$$

Probability propagation

$$\left. \begin{array}{l} \text{Bayesian network} \\ + \\ \text{Evidence: } \mathbf{X}_E = \mathbf{x}_e \end{array} \right\} \Rightarrow P(\mathbf{X}_Q | \mathbf{x}_e)?$$




Example: Calculate $P(X = x | G = g)$

$$P(x, g) = \sum_e P(e) \cdot P(g | e) \sum_l P(l | e) \cdot P(x | l, g)$$
$$P(X = x | G = g) = \frac{P(x, g)}{\sum_{x'} P(X = x', g)} \propto P(x, g)$$

Probability propagation

$$\left. \begin{array}{l} \text{Bayesian network} \\ + \\ \text{Evidence: } \mathbf{X}_E = \mathbf{x}_e \end{array} \right\} \Rightarrow P(\mathbf{X}_Q | \mathbf{x}_e)?$$



Operations to calculate $P(\mathbf{X}_Q | \mathbf{x}_e)$


Restriction: Restrict domain of a potential (e.g., $P(g|E)$ from $P(G|E)$)

Combination: Multiplication of potentials (e.g., $P(l|e) \cdot P(X|l,g)$)

Marginalisation: Sum/integrate out a variable from a potential, e.g., the operation $\sum_l P(l|e) \cdot P(x|l,g)$, which removes L from the potential over $\{L, X, E, G\}$ and results in $P(x|e,g)$ over $\{X, E, G\}$.

Probability propagation

$$\left. \begin{array}{l} \text{Bayesian network} \\ + \\ \text{Evidence: } \mathbf{X}_E = \mathbf{x}_e \end{array} \right\} \Rightarrow P(\mathbf{X}_Q | \mathbf{x}_e)?$$



Requirements for efficient calculation of $P(\mathbf{X}_Q | \mathbf{x}_e)$

- Constraints wrt. structure:
 - Size of combined potentials
- Constraints wrt. distributions:
 - Ability to **perform** operations
 - Ability to **represent results** of operations

Computation of *conditional distributions*

Given $\mathbf{X} = X_1, \dots, X_k \subset \mathcal{V}$, $x_i \in \text{dom}(X_i)$, $\mathbf{Z} = Z_1, \dots, Z_l \subset \mathcal{V}$, compute

$$p(\mathbf{Z} \mid \mathbf{X} = \mathbf{x})$$

Especially: *Single variable posterior distributions*: for each $Z \notin \mathbf{X}$ compute $p(Z \mid \mathbf{X} = \mathbf{x})$.

Problem reduction

This problem can be reduced to the computation of partial distributions, i.e., functions

$$p(\mathbf{Z}, \mathbf{X} = \mathbf{x})$$

(function defined on $\text{dom}(\mathbf{Z})$) because

$$p(\mathbf{Z} \mid \mathbf{X} = \mathbf{x}) = p(\mathbf{Z}, \mathbf{X} = \mathbf{x}) / p(\mathbf{X} = \mathbf{x})$$


Direct approach: Denote $\mathbf{V}(= V_1, \dots, V_m) := \mathcal{V} \setminus (\mathbf{X} \cup \mathbf{Z})$, and let \mathbf{v} range over $\text{dom}(\mathbf{V})$. Then for $\mathbf{z} \in \text{dom}(\mathbf{Z})$:

$$\begin{aligned} p(\mathbf{Z} = \mathbf{z}, \mathbf{X} = \mathbf{x}) &= \sum_{\mathbf{v}} p(\mathbf{Z} = \mathbf{z}, \mathbf{X} = \mathbf{x}, \mathbf{V} = \mathbf{v}) \\ &= \sum_{\mathbf{v}} \prod_i p(Y_i \mid \text{pa}(Y_i))(\mathbf{z}, \mathbf{x}, \mathbf{v}) \\ &= \sum_{v_1 \in \text{dom}(V_1)} \cdots \sum_{v_m \in \text{dom}(V_m)} \prod_i p(Y_i \mid \text{pa}(Y_i))(\mathbf{z}, \mathbf{x}, \mathbf{v}) \end{aligned}$$

“Algorithm”: Sum out the v_j one by one, move factors $p(Y_i \mid \text{pa}(Y_i))(\mathbf{z}, \mathbf{x}, \mathbf{v})$ that do not depend on current v_j (because $V_j \notin \{Y_i\} \cup \text{pa}(Y_i)$) out of the sum.

Advantage: Can be used to compute conditional distributions for arbitrary set \mathbf{Z} of query variables.

Example



	X ₁	
	t	f
	.5	.5

	X ₂		
X ₁	t	f	
t	.7	.3	
f	.1	.9	

	X ₃		
X ₁	t	f	
t	.7	.3	
f	.2	.8	

		X ₄	
X ₂	X ₃	t	f
t	t	.9	.1
t	f	.7	.3
f	t	.8	.2
f	f	.4	.6

$$\begin{aligned}
 p(X_2 = x_2, X_4 = f) &= \sum_{x_1, x_3 \in \{t, f\}} p(X_1 = x_1, X_2 = x_2, X_3 = x_3, X_4 = f) \\
 &= \sum_{x_1, x_3} \left[p(X_1 = x_1) p(X_2 = x_2 \mid X_1 = x_1) p(X_3 = x_3 \mid X_1 = x_1) \right. \\
 &\quad \left. p(X_4 = f \mid X_2 = x_2, X_3 = x_3) \right] \\
 &= \dots
 \end{aligned}$$

$$\begin{aligned}
 & p(X_2 = x_2, X_4 = f) \\
 &= \sum_{x_1} p(X_1 = x_1) p(X_2 = x_2 \mid X_1 = x_1) \left[\sum_{x_3} p(X_3 = x_3 \mid X_1 = x_1) \right. \\
 &\quad \left. p(X_4 = f \mid X_2 = x_2, X_3 = x_3) \right] \\
 &= \sum_{x_1} p(X_1 = x_1) p(X_2 = x_2 \mid X_1 = x_1) F_1(X_1 = x_1, X_2 = x_2) = F_2(X_2 = x_2)
 \end{aligned}$$

where

		X_3	
X_1		t	f
t		.7	.3
f		.2	.8

		X_4	
X_2	X_3	t	f
t	t	.9	.1
t	f	.7	.3
f	t	.8	.2
f	f	.4	.6

 \mapsto

		$F_1(X_1, X_2)$	
x_1	x_2		
t	t	.7·.1 + .3·.3 = .16	
t	f	.7·.2 + .3·.6 = .32	
f	t	.2·.1 + .8·.3 = .26	
f	f	.2·.2 + .8·.6 = .52	

and

	X_1	
	t	f
	.5	.5

		X_2	
X_1		t	f
t		.7	.3
f		.1	.9

x_1	x_2	$F_1(X_1, X_2)$
t	t	.16
\vdots	\vdots	\vdots

 \mapsto

x_2	$F_2(X_2)$
t	...
f	...

- Variable elimination is exponential in maximal number of arguments of factors $p(\dots | \dots)$ resp. $F_j(\dots)$ that appear in the summation process.
- This number depends strongly on the network structure
- ... and can also depend strongly on the order in which we sum out the variables!

Approximate inference using sampling

Problem Structure

Input: Evidence $\mathbf{X} = \mathbf{x}$, random variable Z , value $z \in \text{dom}(Z)$.


Output: Find approximation q for $p := P(Z = z \mid \mathbf{X} = \mathbf{x})$.

Grand plan

- Somehow sample instantiations from the domain \mathbf{Y} ; $\mathbf{X} \cup Z \subseteq \mathbf{Y}$.
- Somehow use the samples $\{\mathbf{y}_1, \dots, \mathbf{y}_N\}$ to find the approximation q .

Observation: can use Bayesian network as random generator that produces full instantiations $\mathbf{Y} = \mathbf{y}$ according to distribution $P(\mathbf{Y})$.

Example:



- Generate random numbers r_1, r_2 uniformly from $[0,1]$.
- Set $A = t$ if $r_1 \leq .2$ and $A = f$ else.
- Depending on the value of A and r_2 set B to t or f .

Generation of one random instantiation: linear in size of network.


- Given a sample $\mathbf{y}_1, \dots, \mathbf{y}_N$ of complete instantiations generated (independently) by the sampling algorithm, approximate $P(\mathbf{X} = \mathbf{x})$ as

$$q^* := \frac{1}{N} |\{i \in 1, \dots, N \mid \mathbf{X} = \mathbf{x} \text{ in } \mathbf{y}_i\}|$$

- Similarly, the sample provides an estimate for $P(Z = z, \mathbf{X} = \mathbf{x})$.
- Put together, we can estimate

$$P(Z = z \mid \mathbf{X} = \mathbf{x}) = P(Z = z, \mathbf{X} = \mathbf{x}) / P(\mathbf{X} = \mathbf{x}).$$

Forward Sampling: Illustration




Approximation for $P(Z = z | \mathbf{X} = \mathbf{x})$:

$$\frac{\# \text{ (Red circle)}}{\# \text{ (Green circle)} \cup \# \text{ (Red circle)}}$$

Problem of forward sampling: samples with $\mathbf{X} \neq \mathbf{x}$ are useless!

Goal: find algorithm that samples according to $P(\mathbf{Z} \mid \mathbf{X} = \mathbf{z})$:



- **Principle:** obtain new sample from previous sample by randomly changing the value of only one selected variable.
- **Notation:** Let $\mathbf{Y} = (\mathbf{Z}, \mathbf{X})$ denote all variables in the domain, where $\mathbf{X} = \mathbf{x}$ is observed.

Gibbs sampling

$\mathbf{z}_0 :=$ arbitrary instantiation of \mathbf{Z} .

$\mathbf{y}_0 := (\mathbf{z}_0, \mathbf{x})$.

$t := 1$.

repeat forever

 choose $Z_k \in \mathbf{Z}$

 set $y_{t,j} := y_{t-1,j}$ for all Y_j except the chosen Z_k .

 generate randomly $z_{t,k}$ according to $P\left(Z_k \mid \mathbf{Y} \setminus \{Z_k\} = \mathbf{y}_t^{\downarrow \mathbf{Y} \setminus \{Z_k\}}\right)$

 Store the sampled value in Z_k 's location in \mathbf{y}_t .

$t := t + 1$.

$$\begin{aligned}
& P(Z_k = z_k \mid \mathbf{Y} \setminus \{Z_k\} = \mathbf{y}_t^{\downarrow \mathbf{Y} \setminus \{Z_k\}}) \\
& \propto P(Z_k = z_k, \mathbf{Y} \setminus \{Z_k\} = \mathbf{y}_t^{\downarrow \mathbf{Y} \setminus \{Z_k\}}) \\
& \propto P(Z_k = z_k \mid \text{pa}(Z_k) = \mathbf{y}_t^{\downarrow \text{pa}(Z_k)}). \\
& \prod_{i: Y_i \in \text{ch}(Z_k)} P(Y_i = y_{t,i} \mid \text{pa}(Y_i) \setminus \{Z_k\} = \mathbf{y}_t^{\downarrow \text{pa}(Y_i) \setminus \{Z_k\}}, Z_k = z_k) \quad (*),
\end{aligned}$$

where \propto means: equals up to a constant that does not depend on z_k .


\rightsquigarrow **Note:** To sample a value we only need to consider the **Markov blanket** for Z_k !

To sample $Z_{t,k}$

- Normalize (*) to obtain $P(Z_k \mid \mathbf{Y} \setminus \{Z_k\} = \mathbf{y}_t^{\downarrow \mathbf{Y} \setminus \{Z_k\}})$
 - $P(Z_k = z_k \mid \mathbf{Y} \setminus \{Z_k\} = \mathbf{y}_t^{\downarrow \mathbf{Y} \setminus \{Z_k\}})$ is called the *full conditional* for Z_k .
- Sample value $z_{t,k}$ according to the resulting distribution

Gibbs Random Walk


The process of Gibbs sampling can be understood as a *random walk* in the space of all instantiations with $\mathbf{X} = \mathbf{x}$:



Reachable in one step: instantiations that differ from current one by value assignment to at most one variable (assume randomized choice of variable Z_k).

Code Task: Exact and approximate inference

In this exercise you should implement a Gibbs sampler for the linear Gaussian model



where the distributions are given as

$$f(x) = \mathcal{N}(x|\mu_x, \sigma_x^2) \quad f(y|x) = \mathcal{N}(y|x, \sigma^2) \quad f(z|x) = \mathcal{N}(z|x, \sigma^2).$$

Start with the partial implementation


```
students_gibbs_sampling.ipynb.
```

Things to try out:


- How does changing the number of samples affect the accuracy of the approximation?
- Try experimenting with different values for the initial parameters.
- Are your results sensitive wrt. the starting position? Why (not)?

A more elaborate example: Factor analysis

- *Factor analysis* is a statistical model used to summarize a high-dimensional observation \mathbf{X} of correlated variables by a smaller set \mathbf{Z} of *factors* that a priori are assumed independent.
- **Example:** \mathbf{X} is a set of scores a subject gets from some intelligence-test, \mathbf{Z} models different types of intelligence (e.g., sense of logics, verbal skills, ...).



Mathematical formulation:




- $\mathbf{Z}_i \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$.
- $X_{i,j} \mid \{\mathbf{z}_i, \mathbf{w}_j, \theta_j\} \sim \mathcal{N}(\mathbf{w}_j^\top \mathbf{z}_i, 1/\theta_j)$.
- **Bayesian setting:** Add \mathbf{W}_j 's and θ as r.v.'s with priors.

Relevant questions given a dataset $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$:


- Learning: $p(\mathbf{w}, \boldsymbol{\theta}, \gamma \mid \mathcal{D})$.
- “Understanding” a new example \mathbf{x}^* : $p(\mathbf{z} \mid \mathbf{X} = \mathbf{x}^*, \mathcal{D})$.
- ...

Unfolded model



FA model “unfolded” for three data instances (X_1, X_2, X_3)

Recall local model



Observations

Inspecting the independence properties of unfolded model we see that the

- number of variables (\mathbf{W} and θ) in “separating factor” are manageable.
- posterior cannot be calculated in closed-form because the priors (assumed a priori independent) are not conjugate. (More on this later.)

↪ Approximate inference required.

Full conditional for $p(\mathbf{w}_j \mid \mathbf{w}_{-j}, \boldsymbol{\theta}, \mathbf{x}, \mathbf{z}, \gamma)$

Let $\mathbf{w}_{-j}, \boldsymbol{\theta}, \mathbf{x}, \mathbf{z}, \gamma$ be a configuration over all variables except \mathbf{w}_j . Then

$$p(\mathbf{w}_j \mid \mathbf{w}_{-j}, \boldsymbol{\theta}, \mathbf{x}, \mathbf{z}, \gamma) \propto p(\mathbf{w}_j \mid \gamma) \prod_{i=1}^N p(x_{ij} \mid \mathbf{w}_j, \mathbf{z}_i, \boldsymbol{\theta}_j)$$

With a bit of pencil pushing we find that:

$$p(\mathbf{w}_j \mid \mathbf{w}_{-j}, \boldsymbol{\theta}, \mathbf{x}, \mathbf{z}, \gamma) = \mathcal{N}(\mathbf{w}_j \mid \boldsymbol{\mu}, \mathbf{Q}^{-1}),$$

where

- $\mathbf{Q} \leftarrow \gamma \mathbf{I} + \theta_j \sum_{i=1}^N \mathbf{z}_i \mathbf{z}_i^\top$
- $\boldsymbol{\mu} \leftarrow \mathbf{Q}^{-1} \theta_j \sum_{i=1}^N x_{ij} \mathbf{z}_i$

Full conditional for $p(\gamma \mid \mathbf{w}, \boldsymbol{\theta}, \mathbf{x}, \mathbf{z})$

$$p(\gamma \mid \mathbf{w}, \boldsymbol{\theta}, \mathbf{x}, \mathbf{z}) \propto p(\gamma) \prod_{j=1}^M p(\mathbf{w}_j \mid \gamma)$$

We find that

$$p(\gamma \mid \mathbf{w}, \boldsymbol{\theta}, \mathbf{x}, \mathbf{z}) = \text{Gamma}(\gamma \mid \text{shape}, \text{rate}),$$

where

- $\text{shape} \leftarrow \text{prior_shape} + \frac{M \cdot D}{2}$
- $\text{rate} \leftarrow \text{prior_rate} + \frac{1}{2} \sum_{j=1}^M \mathbf{w}_j^T \mathbf{w}_j$