

# Variational Inference

## Introduction: Bayesian networks

Helge Langseth and Thomas Dyhre Nielsen\*

Oct. 2018

---

\* Some of the introductory slides are stolen from Manfred Jaeger, Aalborg University.



## Universidad de Almería:

- Antonio Salmeron
- Andrés Masegosa
- Rafael Cabañas de Paz
- Rafael Rumí
- Darío Ramón-Lopez

## Aalborg University:

- Thomas D. Nielsen
- Anders L. Madsen


## NTNU:

- Helge Langseth



**Simultaneous Localization and Mapping:** learn a map of the environment and locate current position

## Example 2: Image Segmentation



(source: <http://pubs.niaaa.nih.gov/publications/arh313/243-246.htm>)


- Divide image into small number of regions representing structurally similar areas

**Given** a collection of texts:



**Goal:** automatically learn semantic descriptors for documents and words, that support document clustering, text understanding, information retrieval ...

Micro-array gene expression data:



Which genes are expressed under which conditions? Which are co-regulated, or functionally dependent?

## Common ground in 4 examples:

- Use a **probabilistic model**, typically learned from data (using **statistical learning techniques**)
- Apply probabilistic inference algorithms to use models for prediction (**classification, regression**), structure analysis (**clustering, segmentation**)

## Advantages of probabilistic/statistical methods:

- Principled quantification of prediction uncertainties
- Robust and principled techniques to deal with incomplete information, missing data.

## **Need: probabilistic models that ...**

- can represent models for high-dimensional state spaces
- support efficient learning and inference techniques

## **Probabilistic Graphical Models ...**

- support a structured specification of high-dimensional distributions in terms of low-dimensional factors
- structured representation can be exploited for efficient learning and inference algorithms (sometimes ...)
- graphical representation gives human-friendly design and description possibilities



## Day 1: Bayesian networks – Definition and inference

- Definition of Bayesian networks: Syntax and semantics
- Exact inference
- Approximate inference using MCMC

## Day 2: Variational inference – Introduction and basis

- Approximate inference through the *Kullback-Leibler divergence*
- *Variational Bayes*
- The *mean-field* approach to Variational Bayes

## Day 3: Variational Bayes – cont'd

- Solving the VB equations
- Introducing Exponential families

## Day 4: Scalable Variational Bayes

- Variational message passing
- Stochastic gradient ascent
- Stochastic variational inference

## Day 5: Current approaches and extensions

- Variational Auto Encoders
- Black Box variational inference
- Probabilistic Programming Languages

## Online repository

Materials for the course, including (almost last-version of) the slides, and description of implementation-tasks are available at

<https://github.com/HelgeLangseth/AlmeriaVB>

HelgeLangseth / AlmeriaVB

Unwatch 2 Star 2 Fork 0

Code Issues 0 Pull requests 0 Projects 0 Wiki Insights Settings

Slides and Jupyter notebooks for VB course in Almeria Edit

Manage topics

7 commits 1 branch 0 releases 1 contributor Apache-2.0

Branch: master New pull request Create new file Upload files Find file Clone or download

File	Commit Message	Time
Python	Added Thomas' MCMC exercise	2 days ago
Slides	Added Thomas' MCMC exercise	2 days ago
LICENSE	Initial commit	4 days ago
README.md	Update README.md	4 days ago

README.md

## AlmeriaVB

This repository contains slides and Jupyter notebooks for VB course in Almeria in Oct 2018.

The slides (in the `slides` folder) are named with the lecture number (1 to 5).

## I will assume that these topics are (fairly well) known:

- Probabilities,  $P(X = x)$ ; conditional probabilities,  $P(X = x \mid Y = y)$ .
- Independence,  $\mathbf{X} \perp\!\!\!\perp \mathbf{Y}$ ; conditional independence  $\mathbf{X} \perp\!\!\!\perp \mathbf{Y} \mid \mathbf{Z}$ .
- “Standard” probability calculus:

**Product:**  $P(x, y) = P(x \mid y) \cdot P(y) = P(x) \cdot P(y \mid x)$ .

**Sum-rule:**  $P(x \vee y) = P(x) + P(y) - P(y \wedge x)$ .

**Total probability:**  $P(Y = y) = \sum_{x'} P(y \mid X = x') \cdot P(X = x')$ .

**Bayes rule:**  $P(x \mid y) = P(x) \cdot P(y \mid x) / P(y)$ .


## I will assume that these topics are (mostly) known:

- Bayesian network syntax and semantics.
- Exact inference in Bayesian networks.
- Approximate inference using MCMC.

## Implementation tasks:


- We will get some implementation tasks as we move along.
- You will be supplied partly running Python-code (in the form of Jupyter notebooks).
- You will need to have **Python 3.x** on your computer (<https://www.python.org/downloads/>), and a set of packages (numpy, scipy, matplotlib, jupyter).

## A simple example: “Explosion”




$$P(E, L, G, X, C)$$

# A simple example: “Explosion”




$$P(E, L, G, X, C)$$

# A simple example: “Explosion”



$$P(E, L, G, X, C)$$

# A simple example: “Explosion”



$$\text{pa}(X) = \{L, G\}$$


$$\text{nd}(X) = \{E, L, G\}$$

$$X \perp\!\!\!\perp E \mid \{L, G\}$$

Other d-sep. rules: Jensen&Nielsen (07)

$$P(E, L, G, X, C)$$

# A simple example: “Explosion”




Other d-sep. rules: Jensen&Nielsen (07)

$$\begin{aligned}
 P(E, L, G, X, C) &= P(E) \cdot P(L \mid E) \cdot P(G \mid E, L) \cdot P(X \mid E, L, G) \cdot P(C \mid E, L, G, X) \\
 &= P(E) \cdot P(L \mid E) \cdot P(G \mid E) \cdot P(X \mid L, G) \cdot P(C \mid X)
 \end{aligned}$$

**Markov properties  $\Leftrightarrow$  Factorisation property**





In the distribution  $P$  defined by the BN the following independence relation holds:

$$P(X_i \mid \text{pa}(X_i), \text{rest}(X_i)) = P(X_i \mid \text{pa}(X_i))$$

“ $X_i$  is independent of its non-descendants given its parents”

## Bayesian network syntax

Let  $X_1, \dots, X_n$  be a collection of random variables. A Bayesian network over  $X_1, \dots, X_n$  consists of

- a directed acyclic graph  $\mathcal{G} = \{\mathcal{V}, \mathcal{E}\}$  whose nodes  $\mathcal{V}$  are the variables  $X_1, \dots, X_n$
- a set of local conditional distributions,  $\mathcal{P} = \{p(X_i \mid \text{pa}(X_i)), X_i \in \mathcal{V}\}$ , where  $\text{pa}(X_i)$  are the parents of  $X_i$  in  $\mathcal{G}$  as defined by the edges  $\mathcal{E}$ .


## Bayesian network semantics

A Bayesian network  $\mathcal{N}$  with nodes  $X_1, \dots, X_n$  defines a joint distribution

$$p_{\mathcal{N}}(X_1, \dots, X_n) = \prod_{i=1}^n p(X_i \mid \text{pa}(X_i))$$

## Probability propagation

$$\left. \begin{array}{l} \text{Bayesian network} \\ + \\ \text{Evidence: } \mathbf{X}_E = \mathbf{x}_e \end{array} \right\} \Rightarrow P(\mathbf{X}_Q | \mathbf{x}_e)?$$




Example: Calculate  $P(X = x | G = g)$

$$\begin{aligned} P(x, g) &= \sum_e \sum_l \sum_c P(E = e, L = l, g, X = x, C = c) \\ &= \sum_e \sum_l \sum_c P(e) \cdot P(l | e) \cdot P(g | e) \cdot P(x | l, g) \cdot P(c | x) \\ &= \sum_e P(e) \cdot P(g | e) \sum_l P(l | e) \cdot P(x | l, g) \sum_c P(c | x) \end{aligned}$$

## Probability propagation

$$\left. \begin{array}{l} \text{Bayesian network} \\ + \\ \text{Evidence: } \mathbf{X}_E = \mathbf{x}_e \end{array} \right\} \Rightarrow P(\mathbf{X}_Q | \mathbf{x}_e)?$$




Example: Calculate  $P(X = x | G = g)$

$$P(x, g) = \sum_e P(e) \cdot P(g | e) \sum_l P(l | e) \cdot P(x | l, g)$$
$$P(X = x | G = g) = \frac{P(x, g)}{\sum_{x'} P(X = x', g)} \propto P(x, g)$$

## Probability propagation

$$\left. \begin{array}{l} \text{Bayesian network} \\ + \\ \text{Evidence: } \mathbf{X}_E = \mathbf{x}_e \end{array} \right\} \Rightarrow P(\mathbf{X}_Q | \mathbf{x}_e)?$$



## Operations to calculate $P(\mathbf{X}_Q | \mathbf{x}_e)$


**Restriction:** Restrict domain of a potential (e.g.,  $P(g|E)$  from  $P(G|E)$ )

**Combination:** Multiplication of potentials (e.g.,  $P(l|e) \cdot P(X|l,g)$ )

**Marginalisation:** Sum/integrate out a variable from a potential, e.g., the operation  $\sum_l P(l|e) \cdot P(x|l,g)$ , which removes  $L$  from the potential over  $\{L, X, E, G\}$  and results in  $P(x|e,g)$  over  $\{X, E, G\}$ .

## Probability propagation

$$\left. \begin{array}{l} \text{Bayesian network} \\ + \\ \text{Evidence: } \mathbf{X}_E = \mathbf{x}_e \end{array} \right\} \Rightarrow P(\mathbf{X}_Q | \mathbf{x}_e)?$$



## Requirements for efficient calculation of $P(\mathbf{X}_Q | \mathbf{x}_e)$

- Constraints wrt. structure:
  - Size of combined potentials
- Constraints wrt. distributions:
  - Ability to **perform** operations
  - Ability to **represent results** of operations

## Bayesian network inference

Inference in the Bayesian network amounts to calculating  $p(\mathbf{Z} = \mathbf{z} \mid \mathbf{X} = \mathbf{x})$ , where

- $\mathbf{X} \subset \mathcal{V}$  are the observed variables, currently taking the configuration  $\mathbf{X} = \mathbf{x}$ .
- $\mathbf{Z} \subseteq \mathcal{V} \setminus \mathbf{X}$  are our variables of interest.

Inference is therefore the tool to answer any probabilistic query we may have in our domain, given a partial (or empty) observation from the domain

- *“What is the chance that Almería will move back to La Liga before 2020, given that they didn’t win during the first four games of this campaign?”*
- *“What is the probability of Google going bankrupt before you are done with your education?”*

## Exact inference



## Computation of *conditional distributions*

Given  $\mathbf{X} = X_1, \dots, X_k \subset \mathcal{V}$ ,  $x_i \in \text{dom}(X_i)$ ,  $\mathbf{Z} = Z_1, \dots, Z_l \subset \mathcal{V}$ , compute

$$p(\mathbf{Z} \mid \mathbf{X} = \mathbf{x})$$

Especially: *Single variable posterior distributions*: for each  $Z \notin \mathbf{X}$  compute  $p(Z \mid \mathbf{X} = \mathbf{x})$ .

## Problem reduction

This problem can be reduced to the computation of partial distributions, i.e., functions

$$p(\mathbf{Z}, \mathbf{X} = \mathbf{x})$$

(function defined on  $\text{dom}(\mathbf{Z})$ ) because

$$p(\mathbf{Z} \mid \mathbf{X} = \mathbf{x}) = p(\mathbf{Z}, \mathbf{X} = \mathbf{x}) / p(\mathbf{X} = \mathbf{x})$$


**Direct approach:** Denote  $\mathbf{V}(= V_1, \dots, V_m) := \mathcal{V} \setminus (\mathbf{X} \cup \mathbf{Z})$ , and let  $\mathbf{v}$  range over  $\text{dom}(\mathbf{V})$ . Then for  $\mathbf{z} \in \text{dom}(\mathbf{Z})$ :

$$\begin{aligned} p(\mathbf{Z} = \mathbf{z}, \mathbf{X} = \mathbf{x}) &= \sum_{\mathbf{v}} p(\mathbf{Z} = \mathbf{z}, \mathbf{X} = \mathbf{x}, \mathbf{V} = \mathbf{v}) \\ &= \sum_{\mathbf{v}} \prod_i p(Y_i \mid \text{pa}(Y_i))(\mathbf{z}, \mathbf{x}, \mathbf{v}) \\ &= \sum_{v_1 \in \text{dom}(V_1)} \cdots \sum_{v_m \in \text{dom}(V_m)} \prod_i p(Y_i \mid \text{pa}(Y_i))(\mathbf{z}, \mathbf{x}, \mathbf{v}) \end{aligned}$$

**“Algorithm”:** Sum out the  $v_j$  one by one, move factors  $p(Y_i \mid \text{pa}(Y_i))(\mathbf{z}, \mathbf{x}, \mathbf{v})$  that do not depend on current  $v_j$  (because  $V_j \notin \{Y_i\} \cup \text{pa}(Y_i)$ ) out of the sum.

**Advantage:** Can be used to compute conditional distributions for arbitrary set  $\mathbf{Z}$  of query variables.

## Example



	X <sub>1</sub>	
	t	f
	.5	.5

	X <sub>2</sub>		
X <sub>1</sub>	t	f	
t	.7	.3	
f	.1	.9	

	X <sub>3</sub>		
X <sub>1</sub>	t	f	
t	.7	.3	
f	.2	.8	

		X <sub>4</sub>	
X <sub>2</sub>	X <sub>3</sub>	t	f
t	t	.9	.1
t	f	.7	.3
f	t	.8	.2
f	f	.4	.6

$$\begin{aligned}
 p(X_2 = x_2, X_4 = f) &= \sum_{x_1, x_3 \in \{t, f\}} p(X_1 = x_1, X_2 = x_2, X_3 = x_3, X_4 = f) \\
 &= \sum_{x_1, x_3} \left[ p(X_1 = x_1) p(X_2 = x_2 \mid X_1 = x_1) p(X_3 = x_3 \mid X_1 = x_1) \right. \\
 &\quad \left. p(X_4 = f \mid X_2 = x_2, X_3 = x_3) \right] \\
 &= \dots
 \end{aligned}$$

$$\begin{aligned}
 & p(X_2 = x_2, X_4 = f) \\
 &= \sum_{x_1} p(X_1 = x_1) p(X_2 = x_2 \mid X_1 = x_1) \left[ \sum_{x_3} p(X_3 = x_3 \mid X_1 = x_1) \right. \\
 &\quad \left. p(X_4 = f \mid X_2 = x_2, X_3 = x_3) \right] \\
 &= \sum_{x_1} p(X_1 = x_1) p(X_2 = x_2 \mid X_1 = x_1) F_1(X_1 = x_1, X_2 = x_2) = F_2(X_2 = x_2)
 \end{aligned}$$

where

		$X_3$	
$X_1$		$t$	$f$
$t$		.7	.3
$f$		.2	.8

		$X_4$	
$X_2$	$X_3$	$t$	$f$
$t$	$t$	.9	.1
$t$	$f$	.7	.3
$f$	$t$	.8	.2
$f$	$f$	.4	.6

 $\mapsto$ 

$x_1$	$x_2$	$F_1(X_1, X_2)$
$t$	$t$	.7·.1 + .3·.3 = .16
$t$	$f$	.7·.2 + .3·.6 = .32
$f$	$t$	.2·.1 + .8·.3 = .26
$f$	$f$	.2·.2 + .8·.6 = .52

and

	$X_1$
	$t$ $f$
	.5 .5

$X_1$	$X_2$	
	$t$ $f$	
$t$	.7 .3	
$f$	.1 .9	

$x_1$	$x_2$	$F_1(X_1, X_2)$
$t$	$t$	.16
$\vdots$	$\vdots$	$\vdots$

 $\mapsto$ 

$x_2$	$F_2(X_2)$
$t$	...
$f$	...

- Variable elimination is exponential in maximal number of arguments of factors  $p(\dots | \dots)$  resp.  $F_j(\dots)$  that appear in the summation process.
- This number depends strongly on the network structure
- ... and can also depend strongly on the order in which we sum out the variables!

## Approximate inference using sampling

## Problem Structure

**Input:** Evidence  $\mathbf{X} = \mathbf{x}$ , random variable  $Z$ , value  $z \in \text{dom}(Z)$ .


**Output:** Find approximation  $q$  for  $p := P(Z = z \mid \mathbf{X} = \mathbf{x})$ .

## Grand plan

- Somehow sample instantiations from the domain  $\mathbf{Y}$ ;  $\mathbf{X} \cup Z \subseteq \mathbf{Y}$ .
- Somehow use the samples  $\{\mathbf{y}_1, \dots, \mathbf{y}_N\}$  to find the approximation  $q$ .

Observation: can use Bayesian network as random generator that produces full instantiations  $\mathbf{Y} = \mathbf{y}$  according to distribution  $P(\mathbf{Y})$ .

## Example:



- Generate random numbers  $r_1, r_2$  uniformly from  $[0,1]$ .
- Set  $A = t$  if  $r_1 \leq .2$  and  $A = f$  else.
- Depending on the value of  $A$  and  $r_2$  set  $B$  to  $t$  or  $f$ .

Generation of one random instantiation: linear in size of network.




- Given a sample  $\mathbf{y}_1, \dots, \mathbf{y}_N$  of complete instantiations generated (independently) by the sampling algorithm, approximate  $P(\mathbf{X} = \mathbf{x})$  as

$$q^* := \frac{1}{N} |\{i \in 1, \dots, N \mid \mathbf{X} = \mathbf{x} \text{ in } \mathbf{y}_i\}|$$

- Similarly, the sample provides an estimate for  $P(Z = z, \mathbf{X} = \mathbf{x})$ .
- Put together, we can estimate

$$P(Z = z \mid \mathbf{X} = \mathbf{x}) = P(Z = z, \mathbf{X} = \mathbf{x}) / P(\mathbf{X} = \mathbf{x}).$$

# Forward Sampling: Illustration




Approximation for  $P(Z = z | \mathbf{X} = \mathbf{x})$ :

$$\frac{\# \text{ (red circle)}}{\# \text{ (green circle)} \cup \text{ (red circle)}}$$

**Problem of forward sampling:** samples with  $\mathbf{X} \neq \mathbf{x}$  are useless!

**Goal:** find algorithm that samples according to  $P(\mathbf{Z} \mid \mathbf{X} = \mathbf{z})$ :



- **Principle:** obtain new sample from previous sample by randomly changing the value of only one selected variable.
- **Notation:** Let  $\mathbf{Y} = (\mathbf{Z}, \mathbf{X})$  denote all variables in the domain, where  $\mathbf{X} = \mathbf{x}$  is observed.

## Gibbs sampling

$\mathbf{z}_0 :=$  arbitrary instantiation of  $\mathbf{Z}$ .

$\mathbf{y}_0 := (\mathbf{z}_0, \mathbf{x})$ .

$t := 1$ .

repeat forever

    choose  $Z_k \in \mathbf{Z}$

    set  $y_{t,j} := y_{t-1,j}$  for all  $Y_j$  except the chosen  $Z_k$ .

    generate randomly  $z_{t,k}$  according to  $P\left(Z_k \mid \mathbf{Y} \setminus \{Z_k\} = \mathbf{y}_t^{\downarrow \mathbf{Y} \setminus \{Z_k\}}\right)$

    Store the sampled value in  $Z_k$ 's location in  $\mathbf{y}_t$ .

$t := t + 1$ .

$$\begin{aligned}
& P(Z_k = z_k \mid \mathbf{Y} \setminus \{Z_k\} = \mathbf{y}_t^{\downarrow \mathbf{Y} \setminus \{Z_k\}}) \\
& \propto P(Z_k = z_k, \mathbf{Y} \setminus \{Z_k\} = \mathbf{y}_t^{\downarrow \mathbf{Y} \setminus \{Z_k\}}) \\
& \propto P(Z_k = z_k \mid \text{pa}(Z_k) = \mathbf{y}_t^{\downarrow \text{pa}(Z_k)}). \\
& \prod_{i: Y_i \in \text{ch}(Z_k)} P(Y_i = y_{t,i} \mid \text{pa}(Y_i) \setminus \{Z_k\} = \mathbf{y}_t^{\downarrow \text{pa}(Y_i) \setminus \{Z_k\}}, Z_k = z_k) \quad (*),
\end{aligned}$$

where  $\propto$  means: equals up to a constant that does not depend on  $z_k$ .


$\rightsquigarrow$  **Note:** To sample a value we only need to consider the **Markov blanket** for  $Z_k$ !

### To sample $Z_{t,k}$

- Normalize (\*) to obtain  $P(Z_k \mid \mathbf{Y} \setminus \{Z_k\} = \mathbf{y}_t^{\downarrow \mathbf{Y} \setminus \{Z_k\}})$ 
  - $P(Z_k = z_k \mid \mathbf{Y} \setminus \{Z_k\} = \mathbf{y}_t^{\downarrow \mathbf{Y} \setminus \{Z_k\}})$  is called the *full conditional* for  $Z_k$ .
- Sample value  $z_{t,k}$  according to the resulting distribution

# Gibbs Random Walk


The process of Gibbs sampling can be understood as a *random walk* in the space of all instantiations with  $\mathbf{X} = \mathbf{x}$ :



Reachable in one step: instantiations that differ from current one by value assignment to at most one variable (assume randomized choice of variable  $Z_k$ ).

**Code Task: Exact and approximate inference**

In this exercise you should implement a Gibbs sampler for the linear Gaussian model



where the distributions are given as

$$f(x) = \mathcal{N}(x|\mu_x, \sigma_x^2) \quad f(y|x) = \mathcal{N}(y|x, \sigma^2) \quad f(z|x) = \mathcal{N}(z|x, \sigma^2).$$

Start with the partial implementation


```
students_gibbs_sampling.ipynb.
```

**Things to try out:**


- How does changing the number of samples affect the accuracy of the approximation?
- Try experimenting with different values for the initial parameters.
- Are your results sensitive wrt. the starting position? Why (not)?

## A more elaborate example: Factor analysis

- *Factor analysis* is a statistical model used to summarize a high-dimensional observation  $\mathbf{X}$  of correlated variables by a smaller set  $\mathbf{Z}$  of *factors* that a priori are assumed independent.
- **Example:**  $\mathbf{X}$  is a set of scores a subject gets from some intelligence-test,  $\mathbf{Z}$  models different types of intelligence (e.g., sense of logics, verbal skills, ...).



### Mathematical formulation:




- $\mathbf{Z}_i \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ .
- $X_{i,j} \mid \{\mathbf{z}_i, \mathbf{w}_j, \theta_j\} \sim \mathcal{N}(\mathbf{w}_j^\top \mathbf{z}_i, 1/\theta_j)$ .
- **Bayesian setting:** Add  $\mathbf{W}_j$ 's and  $\theta$  as r.v.'s with priors.

Relevant questions given a dataset  $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ :

- Learning:  $p(\mathbf{w}, \boldsymbol{\theta}, \gamma \mid \mathcal{D})$ .
- “Understanding” a new example  $\mathbf{x}^*$ :  $p(\mathbf{z} \mid \mathbf{X} = \mathbf{x}^*, \mathcal{D})$ .
- ...




## Unfolded model



FA model “unfolded” for three data instances ( $X_1, X_2, X_3$ )

## Recall local model



## Observations

Inspecting the independence properties of unfolded model we see that the

- number of variables ( $W$  and  $\theta$ ) in “separating factor” are manageable.
- posterior cannot be calculated in closed-form because the priors (assumed a priori independent) are not conjugate. (More on this later.)

↪ Approximate inference required.

## Full conditional for $p(\mathbf{w}_j \mid \mathbf{w}_{-j}, \boldsymbol{\theta}, \mathbf{x}, \mathbf{z}, \gamma)$

Let  $\mathbf{w}_{-j}, \boldsymbol{\theta}, \mathbf{x}, \mathbf{z}, \gamma$  be a configuration over all variables except  $\mathbf{w}_j$ . Then

$$p(\mathbf{w}_j \mid \mathbf{w}_{-j}, \boldsymbol{\theta}, \mathbf{x}, \mathbf{z}, \gamma) \propto p(\mathbf{w}_j \mid \gamma) \prod_{i=1}^N p(x_{ij} \mid \mathbf{w}_j, \mathbf{z}_i, \boldsymbol{\theta}_j)$$

With a bit of pencil pushing we find that:

$$p(\mathbf{w}_j \mid \mathbf{w}_{-j}, \boldsymbol{\theta}, \mathbf{x}, \mathbf{z}, \gamma) = \mathcal{N}(\mathbf{w}_j \mid \boldsymbol{\mu}, \mathbf{Q}^{-1}),$$

where

- $\mathbf{Q} \leftarrow \gamma \mathbf{I} + \theta_j \sum_{i=1}^N \mathbf{z}_i \mathbf{z}_i^\top$
- $\boldsymbol{\mu} \leftarrow \mathbf{Q}^{-1} \theta_j \sum_{i=1}^N x_{ij} \mathbf{z}_i$

**Full conditional for  $p(\gamma \mid \mathbf{w}, \boldsymbol{\theta}, \mathbf{x}, \mathbf{z})$**

$$p(\gamma \mid \mathbf{w}, \boldsymbol{\theta}, \mathbf{x}, \mathbf{z}) \propto p(\gamma) \prod_{j=1}^M p(\mathbf{w}_j \mid \gamma)$$

We find that

$$p(\gamma \mid \mathbf{w}, \boldsymbol{\theta}, \mathbf{x}, \mathbf{z}) = \text{Gamma}(\gamma \mid \text{shape}, \text{rate}),$$

where

- $\text{shape} \leftarrow \text{prior\_shape} + \frac{M \cdot D}{2}$
- $\text{rate} \leftarrow \text{prior\_rate} + \frac{1}{2} \sum_{j=1}^M \mathbf{w}_j^T \mathbf{w}_j$